

Linear regression

This week you will:

- how to prepare data for Spark MLlib
- how to make a predictions by linear regression
- how to estimate prediction quality



```
bike_sharing = spark_session.read.csv(  
    "/user/pmezentsev/regression_bike_sharing/day",  
    header=True)
```

bike_sharing.limit(5).toPandas()

instant		dteday	season	yr	mnth	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	casual	registered	cnt
0	1	2011-01-01	1	0	1	0	6	0	2	0.344167	0.363625	0.805833	0.160446	331	654	985
1	2	2011-01-02	1	0	1	0	0	0	2	0.363478	0.353739	0.696087	0.248539	131	670	801
2	3	2011-01-03	1	0	1	0	1	1	1	0.196364	0.189405	0.437273	0.248309	120	1229	1349
3	4	2011-01-04	1	0	1	0	2	1	1	0.2	0.212122	0.590435	0.160296	108	1454	1562
4	5	2011-01-05	1	0	1	0	3	1	1	0.226957	0.22927	0.436957	0.1869	82	1518	1600

bike_sharing.limit(5).toPandas()



instant		dteday	season	yr	mnth	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	casual	registered	cnt
0	1	2011-01-01	1	0	1	0	6	0	2	0.344167	0.363625	0.805833	0.160446	331	654	985
1	2	2011-01-02	1	0	1	0	0	0	2	0.363478	0.353739	0.696087	0.248539	131	670	801
2	3	2011-01-03	1	0	1	0	1	1	1	0.196364	0.189405	0.437273	0.248309	120	1229	1349
3	4	2011-01-04	1	0	1	0	2	1	1	0.2	0.212122	0.590435	0.160296	108	1454	1562
4	5	2011-01-05	1	0	1	0	3	1	1	0.226957	0.22927	0.436957	0.1869	82	1518	1600




```
bike_sharing.limit(5).toPandas()
```



instant		dteday	season	yr	mnth	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	casual	registered	cnt
0	1	2011-01-01	1	0	1	0	6	0	2	0.344167	0.363625	0.805833	0.160446	331	654	985
1	2	2011-01-02	1	0	1	0	0	0	2	0.363478	0.353739	0.696087	0.248539	131	670	801
2	3	2011-01-03	1	0	1	0	1	1	1	0.196364	0.189405	0.437273	0.248309	120	1229	1349
3	4	2011-01-04	1	0	1	0	2	1	1	0.2	0.212122	0.590435	0.160296	108	1454	1562
4	5	2011-01-05	1	0	1	0	3	1	1	0.226957	0.22927	0.436957	0.1869	82	1518	1600



bike_sharing.limit(5).toPandas()



instant		dteday	season	yr	mnth	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	casual	registered	cnt
0	1	2011-01-01	1	0	1	0	6	0	2	0.344167	0.363625	0.805833	0.160446	331	654	985
1	2	2011-01-02	1	0	1	0	0	0	2	0.363478	0.353739	0.696087	0.248539	131	670	801
2	3	2011-01-03	1	0	1	0	1	1	1	0.196364	0.189405	0.437273	0.248309	120	1229	1349
3	4	2011-01-04	1	0	1	0	2	1	1	0.2	0.212122	0.590435	0.160296	108	1454	1562
4	5	2011-01-05	1	0	1	0	3	1	1	0.226957	0.22927	0.436957	0.1869	82	1518	1600



```
bike_sharing.printSchema()
```

```
root
```

```
| -- instant: string (nullable = true)  
| -- dteday: string (nullable = true)  
| -- season: string (nullable = true)  
| -- yr: string (nullable = true)  
| -- mnth: string (nullable = true)  
| -- holiday: string (nullable = true)  
| -- weekday: string (nullable = true)  
| -- workingday: string (nullable = true)  
| -- weathersit: string (nullable = true)  
| -- temp: string (nullable = true)  
| -- atemp: string (nullable = true)  
| -- hum: string (nullable = true)  
| -- windspeed: string (nullable = true)  
| -- casual: string (nullable = true)  
| -- registered: string (nullable = true)  
| -- cnt: string (nullable = true)
```

```
bike_sharing01 = bike_sharing.select(  
    bike_sharing.season.astype("int"),  
    bike_sharing.yr.astype("int"),  
    bike_sharing.mnth.astype("int"),  
    bike_sharing.holiday.astype("int"),  
    bike_sharing.weekday.astype("int"),  
    bike_sharing.workingday.astype("int"),  
    bike_sharing.weathersit.astype("int"),  
    bike_sharing.temp.astype("double"),  
    bike_sharing.atemp.astype("double"),  
    bike_sharing.hum.astype("double"),  
    bike_sharing.windspeed.astype("double"),  
    bike_sharing.cnt.astype("int").alias("label")  
)
```




```
bike_sharing01 = bike_sharing.select(  
    bike_sharing.season.astype("int"),  
    bike_sharing.yr.astype("int"),  
    bike_sharing.mnth.astype("int"),  
    bike_sharing.holiday.astype("int"),  
    bike_sharing.weekday.astype("int"),  
    bike_sharing.workingday.astype("int"),  
    bike_sharing.weathersit.astype("int"),  
    bike_sharing.temp.astype("double"),  
    bike_sharing.atemp.astype("double"),  
    bike_sharing.hum.astype("double"),  
    bike_sharing.windspeed.astype("double"),  
    bike_sharing.cnt.astype("int").alias("label")  
)
```

```
bike_sharing01 = bike_sharing.select(  
    bike_sharing.season.astype("int"),  
    bike_sharing.yr.astype("int"),  
    bike_sharing.mnth.astype("int"),  
    bike_sharing.holiday.astype("int"),  
    bike_sharing.weekday.astype("int"),  
    bike_sharing.workingday.astype("int"),  
    bike_sharing.weathersit.astype("int"),  
    bike_sharing.temp.astype("double"),  
    bike_sharing.atemp.astype("double"),  
    bike_sharing.hum.astype("double"),  
    bike_sharing.windspeed.astype("double"),  
    bike_sharing.cnt.astype("int").alias("label")  
)
```

```
bike_sharing01.limit(5).toPandas()
```

	season	yr	mnth	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	label
0	1	0	1	0	6	0	2	0.344	0.364	0.806	0.160	985
1	1	0	1	0	0	0	2	0.363	0.354	0.696	0.249	801
2	1	0	1	0	1	1	1	0.196	0.189	0.437	0.248	1349
3	1	0	1	0	2	1	1	0.200	0.212	0.590	0.160	1562
4	1	0	1	0	3	1	1	0.227	0.229	0.437	0.187	1600

Linear regression



Train test splitting



```
train, test = bike_sharing.randomSplit([0.7,0.3])
```



Train



Test



Vector

```
from pyspark.ml.feature import VectorAssembler

assembler = VectorAssembler()\
    .setInputCols(["season",
                    "yr",
                    "mnth",
                    "holiday",
                    "weekday",
                    "weathersit",
                    "temp",
                    "atemp",
                    "hum",
                    "windspeed"])\
    .setOutputCol("features")

train01 = assembler.transform(train)
```



```
train01.limit(5).toPandas()
```

	season	yr	mnth	holiday	weekday	workingday	weathersit	temp	atemp	hum	windspeed	label	features
0	1	0	1	0	0	0	1	0.138	0.116	0.434	0.362	822	[1.0, 0.0, 1.0, 0.0, 0.0, 1.0, 0.138333, 0.116...
1	1	0	1	0	0	0	1	0.232	0.234	0.484	0.188	1204	[1.0, 0.0, 1.0, 0.0, 0.0, 1.0, 0.231667, 0.234...
2	1	0	1	0	0	0	2	0.363	0.354	0.696	0.249	801	[1.0, 0.0, 1.0, 0.0, 0.0, 2.0, 0.363478, 0.353...
3	1	0	1	0	1	1	1	0.097	0.118	0.492	0.158	1416	[1.0, 0.0, 1.0, 0.0, 1.0, 1.0, 0.0973913, 0.11...
4	1	0	1	0	1	1	1	0.151	0.151	0.483	0.223	1321	[1.0, 0.0, 1.0, 0.0, 1.0, 1.0, 0.150833, 0.150...

```
train02 = train01.select("features", "label")  
train02.limit(5).toPandas()
```

	features	label
0	[1.0, 0.0, 1.0, 0.0, 0.0, 1.0, 0.138333, 0.116...	822
1	[1.0, 0.0, 1.0, 0.0, 0.0, 1.0, 0.231667, 0.234...	1204
2	[1.0, 0.0, 1.0, 0.0, 0.0, 2.0, 0.363478, 0.353...	801
3	[1.0, 0.0, 1.0, 0.0, 1.0, 1.0, 0.0973913, 0.11...	1416
4	[1.0, 0.0, 1.0, 0.0, 1.0, 1.0, 0.150833, 0.150...	1321

```
from pyspark.ml.regression import LinearRegression  
lr = LinearRegression()  
model = lr.fit(train03)
```

```
train03 = model.transform(train03)
train03.limit(5).toPandas()
```

	features	label	prediction
0	[1.0, 0.0, 1.0, 0.0, 0.0, 1.0, 0.138333, 0.116...	822	678.100
1	[1.0, 0.0, 1.0, 0.0, 0.0, 1.0, 0.231667, 0.234...	1204	1752.152
2	[1.0, 0.0, 1.0, 0.0, 0.0, 2.0, 0.363478, 0.353...	801	1389.662
3	[1.0, 0.0, 1.0, 0.0, 1.0, 1.0, 0.0973913, 0.11...	1416	1226.827
4	[1.0, 0.0, 1.0, 0.0, 1.0, 1.0, 0.150833, 0.150...	1321	1273.469

```
test01 = assembler.transform(test01)
test02 = test02.select("features", "label")
test03 = model.transform(test03)

test03.limit(3).toPandas()
```

	features	label	prediction
0	[1.0, 0.0, 1.0, 0.0, 0.0, 1.0, 0.0965217, 0.09...	986	883.265574
1	[1.0, 0.0, 1.0, 0.0, 0.0, 2.0, 0.363478, 0.353...	801	1386.370173
2	[1.0, 0.0, 1.0, 0.0, 1.0, 1.0, 0.150833, 0.150...	1321	1265.653788

```
from pyspark.ml.evaluation import RegressionEvaluator  
evaluator = RegressionEvaluator()
```



Coefficient of Determination

R^2


```
evaluator.evaluate(test04,  
                    {evaluator.metricName: "r2"})
```

0.762


0.76




0




0.50



0.80



0.90



0.99

Root mean-square Error

$$\text{RMSE} = \sqrt{\frac{\sum err_i^2}{n}}$$

```
evaluator.evaluate(test04,  
                    {evaluator.metricName: "rmse"})
```

912.341

How often did we miss less than on 300 bikes?


```
test03.select(f.abs(f.col("label")-f.col("prediction")).alias("diff"))\  
  .limit(3).toPandas()
```

	diff
0	102.734426
1	585.370173
2	55.346212

```
test04.select(f.abs(f.col("label")-f.col("prediction")).alias("diff"))\  
    .select(f.when(f.col("diff")<300, 1).otherwise(0).alias("is_accurate"))\  
    .limit(3).toPandas()
```

is_accurate	
0	1
1	0
2	1

```
test03.select(f.abs(f.col("label")-f.col("prediction")).alias("diff"))\  
  .select(f.when(f.col("diff")<300, 1).otherwise(0).alias("is_accurate"))\  
  .agg(f.mean("is_accurate").alias("accuracy"))\  
  .toPandas()
```

accuracy	
0	0.321101

You have learned

- How to prepare data for Spark MLlib
- How to make a predictions by linear regression
- How to estimate prediction quality