


Logistic regression









In this video you will:

- get acquainted with problem of credit risk assessment
- learn how train logistic regression on big data
- learn how to evaluate its performance



```
data = spark_session.read.csv(  
    "/user/pmezentsev/default_of_credit_card_clients",  
    header=True)
```




```
data.printSchema()
```

```
root
```

```
|-- ID: string (nullable = true)
|-- LIMIT_BAL: string (nullable = true)
|-- SEX: string (nullable = true)
|-- EDUCATION: string (nullable = true)
|-- MARRIAGE: string (nullable = true)
|-- AGE: string (nullable = true)
|-- PAY_0: string (nullable = true)
|-- PAY_2: string (nullable = true)
|-- PAY_3: string (nullable = true)
|-- PAY_4: string (nullable = true)
|-- PAY_5: string (nullable = true)
|-- PAY_6: string (nullable = true)
|-- BILL_AMT1: string (nullable = true)
|-- BILL_AMT2: string (nullable = true)
|-- BILL_AMT3: string (nullable = true)
|-- BILL_AMT4: string (nullable = true)
|-- BILL_AMT5: string (nullable = true)
|-- BILL_AMT6: string (nullable = true)
|-- PAY_AMT1: string (nullable = true)
|-- PAY_AMT2: string (nullable = true)
|-- PAY_AMT3: string (nullable = true)
|-- PAY_AMT4: string (nullable = true)
|-- PAY_AMT5: string (nullable = true)
|-- PAY_AMT6: string (nullable = true)
|-- default payment next month: string (nullable = true)
```

ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE
1	20000	2	2	1	24
2	120000	2	2	2	26
3	90000	2	2	2	34
4	50000	2	2	1	37
5	50000	1	2	1	57

Sept.	Aug.	July	June	May	Apr.
PAY_0	PAY_2	PAY_3	PAY_4	PAY_5	PAY_6
2	2	-1	-1	-2	-2
-1	2	0	0	0	2
0	0	0	0	0	0
0	0	0	0	0	0
-1	0	-1	0	0	0



Sept.	Aug.	July	June	May	Apr.
BILL_AMT1	BILL_AMT2	BILL_AMT3	BILL_AMT4	BILL_AMT5	BILL_AMT6
3913	3102	689	0	0	0
2682	1725	2682	3272	3455	3261
29239	14027	13559	14331	14948	15549
46990	48233	49291	28314	28959	29547
8617	5670	35835	20940	19146	19131

ECART

CREDIT INVOICE

ECART THEME
Teststraat 11
1100 AZ Testcity
Telephone: +31(0)201231231
Email: test@test.com
Webshop: http://www.mystore.com

Credit Invoice Date: 13/06/2024

Credit Invoice No: CREDIT-INV2024-002

Order ID: 12

Invoice Address		Shipped to		
test test test test test Noord Holland Netherlands test@test.com 1111111111		test test test test test Noord Holland Netherlands		
Product	Model	Quantity	Unit Price	Total
iPhone	product 11	1	€122.21	€122.21
			Sub-Total:	€101.00
			Pickup from Store:	€0.00
			BTW (21%):	€21.21
			Total:	€122.21

Sept.	Aug.	July	June	May	Apr.
PAY_AMT1	PAY_AMT2	PAY_AMT3	PAY_AMT4	PAY_AMT5	PAY_AMT6
0	689	0	0	0	0
0	1000	1000	1000	0	2000
1518	1500	1000	1000	1000	5000
2000	2019	1200	1100	1059	1000
2000	34481	10000	9000	689	679




default payment next month	
	1
	1
	0
	0
	0


1/0


1/0

Probability


Logistic regression








$$x = \begin{pmatrix} 1 \\ x_1 \\ x_2 \\ \dots \\ x_n \end{pmatrix}$$



$$x = \begin{pmatrix} 1 \\ x_1 \\ x_2 \\ \dots \\ x_n \end{pmatrix}$$


$$y \in (0,1)$$




$$x = \begin{pmatrix} 1 \\ x_1 \\ x_2 \\ \dots \\ x_n \end{pmatrix}$$

$$y \in (0,1)$$

$$w = \begin{pmatrix} w_0 \\ w_1 \\ w_2 \\ \dots \\ w_n \end{pmatrix}$$





$$x = \begin{pmatrix} 1 \\ x_1 \\ x_2 \\ \dots \\ x_n \end{pmatrix}$$

$$y \in (0,1)$$


$$w = \begin{pmatrix} w_0 \\ w_1 \\ w_2 \\ \dots \\ w_n \end{pmatrix}$$

$$x = \begin{pmatrix} 1 \\ x_1 \\ x_2 \\ \dots \\ x_n \end{pmatrix}$$

$$y \in (0,1)$$

$$w = \begin{pmatrix} w_0 \\ w_1 \\ w_2 \\ \dots \\ w_n \end{pmatrix}$$

$$\hat{y} = \frac{1}{1 + e^{-w^T x}}$$





$$\hat{y} = \frac{1}{1 + e^{-w^T x}}$$

$$w = \begin{pmatrix} w_0 \\ w_1 \\ w_2 \\ \dots \\ w_n \end{pmatrix}$$

$$y \in (0, 1)$$

$$x = \begin{pmatrix} 1 \\ x_1 \\ x_2 \\ \dots \\ x_n \end{pmatrix}$$





L-BFGS



```
from pyspark.ml.feature import SQLTransformer
```

```
my_sql_transformer = SQLTransformer(  
    statement="""  
    SELECT  
        cast(LIMIT_BAL          as int),  
        cast(SEX                as int),  
        cast(EDUCATION          as int),  
        cast(MARRIAGE           as int),  
        cast(AGE                as int),  
        cast(PAY_0              as int),  
        cast(PAY_2              as int),  
        cast(PAY_3              as int),  
        cast(PAY_4              as int),  
        cast(PAY_5              as int),  
        cast(PAY_6              as int),  
        cast(BILL_AMT1          as int),  
        cast(BILL_AMT2          as int),  
        cast(BILL_AMT3          as int),  
        cast(BILL_AMT4          as int),  
        cast(BILL_AMT5          as int),  
        cast(BILL_AMT6          as int),  
        cast(PAY_AMT1           as int),  
        cast(PAY_AMT2           as int),  
        cast(PAY_AMT3           as int),  
        cast(PAY_AMT4           as int),  
        cast(PAY_AMT5           as int),  
        cast(PAY_AMT6           as int),  
        cast('default payment next month' as int) as label  
    FROM __THIS__  
    """)
```

```
assembler = VectorAssembler()\
    .setInputCols([\
        'LIMIT_BAL',\
        'SEX',\
        'EDUCATION',\
        'MARRIAGE',\
        'AGE',\
        'PAY_0', 'PAY_2', 'PAY_3', 'PAY_4', 'PAY_5', 'PAY_6',\
        'BILL_AMT1', 'BILL_AMT2', 'BILL_AMT3', 'BILL_AMT4', 'BILL_AMT5', 'BILL_AMT6',\
        'PAY_AMT1', 'PAY_AMT2', 'PAY_AMT3', 'PAY_AMT4', 'PAY_AMT5', 'PAY_AMT6'\
    ])\
    .setOutputCol("features")
```

```
data02 = assembler.transform(data01).select('features', 'label')
```

```
train, test = data02.randomSplit([0.7,0.3], seed=1)
```



```
lr = LogisticRegression()  
lr_model = lr.fit(train)
```

```
scored_test = lr_model.transform(test)
scored_train = lr_model.transform(train)
```

```
scored_test.limit(5).toPandas()
```

	features	label	rawPrediction	probability	pre- diction
0	(10000.0, 2.0, 1.0, 2.0, 23.0, 1.0, -2.0, -2.0...	0	[1.04333251803, -1.04333251803]	[0.739492505681, 0.260507494319]	0.0
1	(20000.0, 1.0, 2.0, 2.0, 23.0, 1.0, -2.0, -2.0...	0	[1.04244635215, -1.04244635215]	[0.739321755498, 0.260678244502]	0.0
2	(20000.0, 1.0, 3.0, 2.0, 40.0, 1.0, -2.0, -2.0...	0	[1.0287134211, -1.0287134211]	[0.736666389661, 0.263333610339]	0.0
3	(20000.0, 2.0, 2.0, 1.0, 26.0, 1.0, -2.0, -2.0...	1	[0.976449927859, -0.976449927859]	[0.726403236781, 0.273596763219]	0.0
4	(20000.0, 2.0, 2.0, 1.0, 48.0, 1.0, -2.0, -2.0...	1	[0.840721980366, -0.840721980366]	[0.698617251621, 0.301382748379]	0.0

```
scored_test.limit(5).toPandas()
```

	features	label	rawPrediction	probability	pre-diction
0	(10000.0, 2.0, 1.0, 2.0, 23.0, 1.0, -2.0, -2.0...	0	[1.04333251803, -1.04333251803]	[0.739492505681, 0.260507494319]	0.0
1	(20000.0, 1.0, 2.0, 2.0, 23.0, 1.0, -2.0, -2.0...	0	[1.04244635215, -1.04244635215]	[0.739321755498, 0.260678244502]	0.0
2	(20000.0, 1.0, 3.0, 2.0, 40.0, 1.0, -2.0, -2.0...	0	[1.0287134211, -1.0287134211]	[0.736666389661, 0.263333610339]	0.0
3	(20000.0, 2.0, 2.0, 1.0, 26.0, 1.0, -2.0, -2.0...	1	[0.976449927859, -0.976449927859]	[0.726403236781, 0.273596763219]	0.0
4	(20000.0, 2.0, 2.0, 1.0, 48.0, 1.0, -2.0, -2.0...	1	[0.840721980366, -0.840721980366]	[0.698617251621, 0.301382748379]	0.0

$$w^T x$$

```
scored_test.limit(5).toPandas()
```

	features	label	rawPrediction	probability	pre-diction
0	(10000.0, 2.0, 1.0, 2.0, 23.0, 1.0, -2.0, -2.0...	0	[1.04333251803, -1.04333251803]	[0.739492505681, 0.260507494319]	0.0
1	(20000.0, 1.0, 2.0, 2.0, 23.0, 1.0, -2.0, -2.0...	0	[1.04244635215, -1.04244635215]	[0.739321755498, 0.260678244502]	0.0
2	(20000.0, 1.0, 3.0, 2.0, 40.0, 1.0, -2.0, -2.0...	0	[1.0287134211, -1.0287134211]	[0.736666389661, 0.263333610339]	0.0
3	(20000.0, 2.0, 2.0, 1.0, 26.0, 1.0, -2.0, -2.0...	1	[0.976449927859, -0.976449927859]	[0.726403236781, 0.273596763219]	0.0
4	(20000.0, 2.0, 2.0, 1.0, 48.0, 1.0, -2.0, -2.0...	1	[0.840721980366, -0.840721980366]	[0.698617251621, 0.301382748379]	0.0

$$w^T x$$

$$\hat{y} = \frac{1}{1 + e^{-w^T x}}$$

```
scored_test.limit(5).toPandas()
```

	features	label	rawPrediction	probability	pre-diction
0	(10000.0, 2.0, 1.0, 2.0, 23.0, 1.0, -2.0, -2.0...	0	[1.04333251803, -1.04333251803]	[0.739492505681, 0.260507494319]	0.0
1	(20000.0, 1.0, 2.0, 2.0, 23.0, 1.0, -2.0, -2.0...	0	[1.04244635215, -1.04244635215]	[0.739321755498, 0.260678244502]	0.0
2	(20000.0, 1.0, 3.0, 2.0, 40.0, 1.0, -2.0, -2.0...	0	[1.0287134211, -1.0287134211]	[0.736666389661, 0.263333610339]	0.0
3	(20000.0, 2.0, 2.0, 1.0, 26.0, 1.0, -2.0, -2.0...	1	[0.976449927859, -0.976449927859]	[0.726403236781, 0.273596763219]	0.0
4	(20000.0, 2.0, 2.0, 1.0, 48.0, 1.0, -2.0, -2.0...	1	[0.840721980366, -0.840721980366]	[0.698617251621, 0.301382748379]	0.0

$$w^T x$$

$$\hat{y} = \frac{1}{1 + e^{-w^T x}}$$

$$\hat{y} > 0,5$$

```
from pyspark.ml.evaluation import BinaryClassificationEvaluator  
evaluator = BinaryClassificationEvaluator()
```


```
evaluator.evaluate(scoped_train, {evaluator.metricName: 'areaUnderROC'})
```

0.722522648109

```
evaluator.evaluate(scoped_test, {evaluator.metricName: 'areaUnderROC'})
```

0.728012080672

ROC



In this video you:

- have got acquainted with problem of credit risk assessment
- have learned how to train logistic regression on big data
- have learned how to evaluate its performance